RUM 2 (RNA-Seq Unified Mapper)

Mike DeLaurentis

University of Pennsylvania

August 15, 2012

Agenda

1. Intro

2. The RUM pipeline

3. Tour of the new RUM
   - Installation
   - Command-line interface
   - Job state management
   - Work distribution

4. Demo
   - Web resources
   - Run tiny job

Agenda

Who am I?

- Working at Penn (ITMAT) since January
- Software engineering background
- Experience in a variety of languages, (Perl, Java, Clojure (lisp), Python, Ruby)
- ... and applications

Agenda

What is RUM?

- *"RUM is an alignment, junction calling, and feature quantification pipeline specifically designed for Illumina RNA-Seq data"*
- Written by Gregory Grant (ggrant@grant.org)
- Runs on Linux / UNIX / Mac
- Distributed (optionally)

What is RUM?

Inputs
- RNA-Seq reads: FASTA or FASTQ, paired or single
- Organism-specific index

What is RUM?

Inputs
- RNA-Seq reads: FASTA or FASTQ, paired or single
- Organism-specific index

Outputs
- Unique and non-unique alignments
- Coverage plots
- Feature quantifications
- Junction calls
- List of novel inferred internal exons

RUM has three distinct phases:

- Preprocessing
- Processing
- Postprocessing

Phase 1: Preprocessing

- Perform some quality checks on reads
- Split reads into N chunks
- Allows reads to be processed by N nodes on a cluster

Phase 2: Processing

- Align all reads against genome using Bowtie
- Align all reads against transcriptome using Bowtie
- Merge genome and transcriptome alignments and identify unmapped reads
- Align unmapped reads against genome using BLAT
- *"This leverages the advantages of both genome and transcriptome mapping as well as combining the speed of Bowtie with the sensitivity and flexibility of Blat."*
- Merge Bowtie and Blat alignments

Phase 3: Postprocessing

- Merge alignments for all chunks together
- Produce coverage plots, junction files
- Find novel internal exons
- Generate some reports

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Agenda

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

RUM 2 Goals

- More code reuse for maintainability
- Improve inter-process communication
- More robust job state management
- Automated tests

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Enhancements in RUM 2

- Standard installation process
- New command-line interface
- Get status of running job
- Restart a job where it left off
- More reliable kill command
- Run a chunk or postprocessing by itself
- Relocatable indexes
- SAM file is closer to conforming to standard

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Installing RUM

- Uses standard Perl Makefile.PL
- Should be familiar to system administrators
- Download tarball from
  https://github.com/PGFI/rum/downloads
- `perl Makefile.PL`
- `make install` (optional)
- Then install indexes...

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Installing Indexes

- RUM needs an index for each organism you want to align against
- Index includes genome, gene annotations, and binary index files for Bowtie
- Index installation is now separate from code installation / upgrade
- Run rum_indexes; it will guide you through the process

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Command-line interface

Usage is rum_runner ACTION [OPTIONS] where action is one of:

- align - Run an alignment
- status - Check the status of a job
- stop - Stop a job (can be restarted later)
- kill - Stop a job and clean it up (to restart from scratch)
- clean - Remove output files for a job
- help - Get help
- version - Show version number

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Running an alignment

Use "rum_runner align" to run an alignment:

```
rum_runner align \
  --output ~/sample123/results \
  --index  ~/rum_indexes/hg19 \
  --name   TestJob \
  --chunks 25 \
  ~/sample123/forward.fq ~/sample123/reverse.fq
```

Intro   Installation
The RUM pipeline   Command-line interface
Tour of the new RUM   Job state management
Demo   Work distribution

Job status

- Use rum_runner status to check on the status of a running job.

```
$ rum_runner status -o ~/sample123/results

Processing in 25 chunks
-----------------------
XXXXXXXXXXXXXXXXXXXXXXXXX Run bowtie on genome
XXXXXXXXXXXXXXXXXXXXXXXXX Parse genome Bowtie output
X XXX  XX XX XXXX  XXXXX  Run bowtie on transcriptome
X XXX  XX XX XXXX  XXXXX  Parse transcriptome Bowtie output
X XXX  XX XX XXXX  XXXXX  Merge unique mappers together
X XXX  XX XX XXXX  XXXXX  Merge non-unique mappers together
X XXX  XX XX XXXX  XXXXX  Make unmapped reads file for blat
X XXX  XX XX XXXX  XXXXX  Run blat on unmapped reads
X XXX  XX XX XXXX  XXXXX  Run mdust on unmapped reads
X XXX  XX XX XXXX  XXXXX  Parse blat output
X XXX  XX XX XXXX  XXXXX  Merge bowtie and blat results
X XX   XX XX  XXX   XXXX  Clean up RUM files
X XX   XX XX  XXX   XXXX  Produce RUM_Unique
X XX   XX XX  XXX   XXXX  Sort RUM_Unique by location
X X    XX XX  XXX   XXXX  Sort cleaned non-unique mappers by ID
X X    XX XX  XXX   XXXX  Remove duplicates from NU
X X    XX XX  XXX   XXXX  Create SAM file
X X    XX XX  XXX   XXXX  Create non-unique stats
X X    XX XX  XXX   XXXX  Sort RUM_NU
X X    XX XX  XXX   XXXX  Generate quants
```

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Job status

```
Postprocessing
--------------
X Merge RUM_NU files
X Make non-unique coverage
X Merge RUM_Unique files
X Compute mapping statistics
X Make unique coverage
X Finish mapping stats
X Merge SAM headers
X Concatenate SAM files
X Merge quants
  make_junctions
  Sort junctions (all, bed) by location
  Sort junctions (all, rum) by location
  Sort junctions (high-quality, bed) by location
  Get inferred internal exons
  Quantify novel exons

All the chunk error log files are empty. That's good.
Main error log file is empty. That's good.

RUM is running (job ids 815718, 815720).
```

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Job state management

- Model the workflow as a state machine
- A completed step transitions the job from one state to another
- Stitch steps together into a workflow
- Determine the state of a job by looking at which output files exist
- Similar to GNU Make (but simpler)
- Basis for a lot of additional features

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Recovering from errors

- In case of failure...
- RUM 2 allows easier recovery
- Running "rum_runner align" again, RUM will determine what state the job was in when it failed
- Just resumes at the next uncompleted step
- Can save *a lot* of time when recovering from infrastructure failure

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Killing a job

- To stop a job and remove all of its output:

    ```
    rum_runner kill -o dir
    ```

- Useful if you've run a job with incorrect parameters and need to start over

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Work distribution

TM & © 2006 The Jim Henson Company

- Automatic support for one multi-core machine (will run each chunk in a separate process by default)
- Built-in support for Sun Grid Engine, with --qsub option
- Easily extensible for other platforms

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Work distribution

For other platforms, you can either extend a Perl class to provide
support, or set up some scripts to run the parts of the job on
different machines:

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Work distribution

For other platforms, you can either extend a Perl class to provide
support, or set up some scripts to run the parts of the job on
different machines:

- Run preprocessing alone:

  ```
  rum_runner -o <dir> --preprocess
  ```

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Work distribution

For other platforms, you can either extend a Perl class to provide
support, or set up some scripts to run the parts of the job on
different machines:

- Run preprocessing alone:

  ```
  rum_runner -o <dir> --preprocess
  ```

- Run chunks one at a time:

  ```
  rum_runner -o <dir> --process --chunk 8
  ```

Intro
The RUM pipeline
Tour of the new RUM
Demo

Installation
Command-line interface
Job state management
Work distribution

Work distribution

For other platforms, you can either extend a Perl class to provide support, or set up some scripts to run the parts of the job on different machines:

- Run preprocessing alone:

  ```
  rum_runner -o <dir> --preprocess
  ```

- Run chunks one at a time:

  ```
  rum_runner -o <dir> --process --chunk 8
  ```

- Run postprocessing alone:

  ```
  rum_runner -o <dir> --postprocess
  ```

Intro
The RUM pipeline
Tour of the new RUM
Demo

Web resources
Run tiny job

Agenda

Intro
The RUM pipeline
Tour of the new RUM
Demo

Web resources
Run tiny job

Web resources

Main github page `https://github.com/PGFI/rum`

User guide `https://github.com/PGFI/rum/wiki`

Issues `https://github.com/PGFI/rum/issues`

Downloads `https://github.com/PGFI/rum/downloads`

Intro
The RUM pipeline
Tour of the new RUM
Demo

Web resources
Run tiny job

To the command line...